



After kernel finishes booting - create first process

Then subsequent processes will be forked

↳ clone process - exact copy
→ both then return from fork

The return values will be different

- Parent will get process ID and child will return 0

waitpid - wait for one process to finish

↳ can capture the exit status of a process

i.e. `child_exit = waitpid(child_pid);`

if the child exits before we call waitpid - we immediately return the exit status

Therefore, it always works!

We will need to implement the following system calls

open, write, read

```
char buf[1];  
fd=open("FILE", O_RDWR | O_CREATE);  
rc=write(fd, "HELLO\n", 6); if(rc < 0) {  
rc=write(fd, "WORLD\n", 6); } failed write  
lseek(fd, 0, SEEK_SET); // go to position zero  
read(fd, buf, 7) // copy from 0 to 7
```

HELLO\nWORLD\n

0 5 6

```
fd2=open("FILE", O_RDWR); // will have a different handle  
and offset
```

```
fd=open("FILE" ....)  
rc=FORK()  
if(rc == 0) {  
    child — read(fd)  
} else {  
    parent read(fd)  
}
```

For each process

- keep track of open files
 - file handle given to app
 - current file position
 - anything else...

File handles 0,1,2

- automatically available in each process

After fork()

- child has same open file handle as parent
- parent and child **share** current file position for each file at time of fork()
- this property is desired when you want a parent and their child to cooperate

After execv

- open files unchanged